



Reading covid19.csv

```
path = "/FileStore/tables/"
fc = path + "covid19.csv"

dfc1 = spark.read.csv( fc, header=True, inferSchema=True )
dfc1.printSchema()
dfc1.show(2)
```

```
root
 |-- Province/State: string (nullable = true)
 |-- Country/Region: string (nullable = true)
 |-- Lat: double (nullable = true)
 |-- Long: double (nullable = true)
 |-- Date: string (nullable = true)
 |-- Confirmed: integer (nullable = true)
 |-- Deaths: integer (nullable = true)
 |-- Recovered: integer (nullable = true)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|Province/State|Country/Region|    Lat|    Long|    Date|Confirmed|Deaths|Recovered|
+-----+-----+-----+-----+-----+-----+-----+
|          null|    Afghanistan|   33.0|   65.0|1/22/20|         0|      0|         0|
|          null|        Albania|41.1533|20.1683|1/22/20|         0|      0|         0|
+-----+-----+-----+-----+-----+-----+-----+
```

only showing top 2 rows

```
print( dfc1.rdd.getNumPartitions() )
print( sc.defaultParallelism )
dfc1 = dfc1.repartition( 24 )
print( dfc1.rdd.getNumPartitions() )
```

```
1
8
24
```

Schema modification, aggregation on states for each

```

from pyspark.sql import functions as fun
dfc2 = dfc1.select( dfc1["Country/Region"].alias("country"),
                    fun.to_timestamp( "Date", "MM/dd/yy" ).alias("date"),
                    dfc1.Confirmed.alias("confirmed"),
                    dfc1.Deaths.alias("deaths"),
                    dfc1.Recovered.alias("recovered"))\
    .groupby("country", "date")\
    .agg( fun.sum("confirmed").alias("confirmed"),
          fun.sum("deaths").alias("deaths"),
          fun.sum("recovered").alias("recovered"))

```

```
dfc2.filter( dfc2.date > fun.lit("2020-04-12") ).show(5)
```

```

+-----+-----+-----+-----+-----+
|      country|      date|confirmed|deaths|recovered|
+-----+-----+-----+-----+-----+
|United Kingdom|2020-05-17 00:00:00|    244995|   34716|    1058|
|      Zambia|2020-05-10 00:00:00|      267|      7|     117|
|      Kosovo|2020-05-04 00:00:00|      855|     26|     403|
|    Bulgaria|2020-04-12 00:00:00|      675|     29|      68|
|      Egypt|2020-04-29 00:00:00|     5268|     380|    1335|
+-----+-----+-----+-----+-----+

```

only showing top 5 rows


Computing the daily statistical changes

```

# Adding newConfirmed ...
from pyspark.sql import Window
win = Window.partitionBy("country").orderBy("date")
dfc3 = dfc2.withColumn( "newConfirmed", dfc2.confirmed - fun.lag( dfc2.confirmed, 1
).over(win) )\
    .withColumn( "newDeaths", dfc2.deaths - fun.lag( dfc2.deaths, 1
).over(win) )\
    .withColumn( "newRecovered", dfc2.recovered - fun.lag( dfc2.recovered, 1
).over(win) )

display( dfc3.filter( (dfc3.country=='Brazil') & (dfc3.newConfirmed > 0) ) )
# >"2020-03-15"

```



2-week binning over newConfirmed

```
# date binning
dfc4 = dfc3.groupby("country",
                    fun.window("date", "14 days").alias("timeInterval") )\
        .agg( fun.sum("newConfirmed").alias("bnc") )
```

```
dfc4.printSchema()
dfc4.filter(dfc4.country=='United Kingdom').show(10, False)
```

```
root
 |-- country: string (nullable = true)
 |-- timeInterval: struct (nullable = false)
 |   |-- start: timestamp (nullable = true)
 |   |-- end: timestamp (nullable = true)
 |-- bnc: long (nullable = true)
```

```
+-----+-----+-----+-----+
|country|timeInterval|bnc|
```


```

+-----+-----+-----+
|United Kingdom|[2020-01-09 00:00:00, 2020-01-23 00:00:00]|null|
|United Kingdom|[2020-01-23 00:00:00, 2020-02-06 00:00:00]|2|
|United Kingdom|[2020-02-06 00:00:00, 2020-02-20 00:00:00]|7|
|United Kingdom|[2020-02-20 00:00:00, 2020-03-05 00:00:00]|77|
|United Kingdom|[2020-03-05 00:00:00, 2020-03-19 00:00:00]|2556|
|United Kingdom|[2020-03-19 00:00:00, 2020-04-02 00:00:00]|27223|
|United Kingdom|[2020-04-02 00:00:00, 2020-04-16 00:00:00]|69618|
|United Kingdom|[2020-04-16 00:00:00, 2020-04-30 00:00:00]|66958|
|United Kingdom|[2020-04-30 00:00:00, 2020-05-14 00:00:00]|64544|
|United Kingdom|[2020-05-14 00:00:00, 2020-05-28 00:00:00]|19153|
+-----+-----+-----+


```

```
dfc5 = dfc4.select("country", dfc4.timeInterval.start.cast("date").alias("date"),
"bnc")
```

```
display( dfc5.filter( (dfc5.country=='Brazil') ) )
```



```
display( dfc5.filter( (dfc5.country=='Norway') ) )
```



```
dfc5.printSchema()
```

```
root
```

```
|-- country: string (nullable = true)
|-- date: date (nullable = true)
|-- bnc: long (nullable = true)
```

```
# mean and std per country for binned new confirmed
```

```
dfc5.filter( "date > '2020-03-15' " )\
  .groupby("country")\
  .agg( fun.round( fun.avg("bnc"), 0 ).alias("mean_bnc"),
        fun.round( fun.stddev("bnc"), 0 ).alias("std_bnc") )\
  .orderBy("mean_bnc", ascending=False)\
  .show()
```

```

+-----+-----+-----+
|      country|mean_bnc| std_bnc|
+-----+-----+-----+
|           US|304157.0|125969.0|
|        Russia| 59959.0| 54435.0|
|        Brazil| 54303.0| 42662.0|
|United Kingdom| 49499.0| 24254.0|
|          Spain| 43625.0| 37188.0|
|          Italy| 38197.0| 27970.0|
|          France| 34362.0| 29331.0|
|         Germany| 33090.0| 27136.0|
|          Turkey| 30303.0| 19894.0|
|           Iran| 21448.0|  7840.0|
|          India| 21264.0| 16669.0|
|          Peru| 19868.0| 15556.0|
|         Canada| 15967.0|  7694.0|
| Saudi Arabia| 11937.0|  8989.0|
|         Belgium| 10861.0|  6987.0|
|          Mexico| 10846.0|  8336.0|
|          Chile|  9868.0|  7127.0|
|        Pakistan|  8733.0|  6889.0|
| Netherlands|  8477.0|  5577.0|
|          Qatar|  7031.0|  5412.0|
+-----+-----+-----+

```

only showing top 20 rows

median per country for binned new confirmed

from pyspark.sql **import** Window

win2 = Window.partitionBy("country")

dfc5.filter(" date > '2020-03-15' ")\

.withColumn("medianBnc", fun.expr("percentile_approx(bnc, .5)").over(win2)

)\

.show()

```

+-----+-----+-----+-----+
| country|      date|   bnc|medianBnc|
+-----+-----+-----+-----+
|    Chad|2020-03-19|    7|      29|
|    Chad|2020-04-02|   16|      29|
|    Chad|2020-04-16|   29|      29|
|    Chad|2020-04-30|  320|      29|
|    Chad|2020-05-14|  173|      29|
|Paraguay|2020-03-19|   58|      89|
|Paraguay|2020-04-02|   92|      89|
|Paraguay|2020-04-16|   78|      89|
|Paraguay|2020-04-30|  501|      89|
|Paraguay|2020-05-14|   89|      89|

```

```

| Russia|2020-03-19| 2630| 57670|
| Russia|2020-04-02| 21713| 57670|
| Russia|2020-04-16| 74909| 57670|
| Russia|2020-04-30|142872| 57670|
| Russia|2020-05-14| 57670| 57670|
| Yemen|2020-03-19| 0| 5|
| Yemen|2020-04-02| 1| 5|
| Yemen|2020-04-16| 5| 5|
| Yemen|2020-04-30| 64| 5|
| Yemen|2020-05-14| 97| 5|
+-----+-----+-----+-----+

```

only showing top 20 rows

```
dfc3.printSchema()
```

```
root
```

```

|-- country: string (nullable = true)
|-- date: timestamp (nullable = true)
|-- confirmed: long (nullable = true)
|-- deaths: long (nullable = true)
|-- recovered: long (nullable = true)
|-- newConfirmed: long (nullable = true)
|-- newDeaths: long (nullable = true)
|-- newRecovered: long (nullable = true)

```

```
# keeping the list of aggregation on 14 days
```

```

dfc6 = dfc3.groupby( "country", fun.window( "date", "14 days" ).start.alias("date")
)\
    .agg( fun.sum("newConfirmed").alias("bnc"), fun.collect_list( "newConfirmed"
).alias("listBnc") )

```

```
dfc6.filter( "country='France'" ).show(10, False)
```

```
dfc3.filter( "country='France'" ).show(10, False)
```

```

+-----+-----+-----+-----+
-----+
|country|date                |bnc  |listBnc
|
+-----+-----+-----+-----+
-----+
|France |2020-01-09 00:00:00|null |[]
|
|France |2020-01-23 00:00:00|6     |[0, 2, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0]
|

```

```
|France |2020-02-06 00:00:00|6      |[0, 0, 5, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
|
|France |2020-02-20 00:00:00|276   |[0, 0, 0, 0, 0, 2, 4, 20, 19, 43, 30, 61, 13, 8
4]
|
|France |2020-03-05 00:00:00|8836 |[92, 276, 303, 177, 83, 575, 499, 0, 1388, 815,
36, 2151, 1032, 1409]
|
|France |2020-03-19 00:00:00|48625|[1846, 1788, 1705, 1780, 3880, 2499, 2978, 3951
, 3851, 4703, 2603, 4462, 7657, 4922] |
|France |2020-04-02 00:00:00|75836|[2180, 5273, 4298, 1912, 3931, 3820, 3894, 4309
, 4372, 3125, 26849, 3682, 4971, 3220]|
```

```
# writting a udf to compute line slop for 14-day list (listBnc)
```

```
import numpy as np
```

```
from sklearn.linear_model import LinearRegression
```

```
@udf( "float" )
```

```
def computeSlope( ll ) :
```

```
    try :
```

```
        zcount = len( [ 0 for i in ll if i==0 ] )
```

```
        # return none if there are not enough data points to fit
```

```
        if zcount > int(len(ll)*1./2.) : return None
```

```
        xx, yy = [], []
```

```
        for i in range( len(ll) ) :
```

```
            if ll[i] > 0 :
```

```
                xx.append( i )
```

```
                yy.append( ll[i] )
```

```
        X = np.array( xx ).reshape(-1,1) # time
```

```
        y = np.array( yy, ) # newConfirmed
```

```
        reg = LinearRegression()
```

```
        reg.fit(X, y)
```

```
        s = reg.score(X,y)
```

```
        # to filter the fit quality (0 : no filter)
```

```
        if s > .0 :
```

```
            # return float( reg.coef_[0] )
```

```
            # line slop computation
```

```
            return float( round( np.arctan(reg.coef_[0])*180./np.pi, 1 ) )
```

```
            # you should call the float function to convert from numpy-float to Python-
```

```
float
```

```
    else : return None
```

```
except ValueError :
```

```
    return None
```



```
dfc7 = dfc6.withColumn( "slope", computeSlope( "listBnc" ) )
```

```
# Printng for brazil
```

```
dfc7.filter("country='Brazil']").show(100, False)
```

```
+-----+-----+-----+-----+
|country|date                |bnc  |listBnc
|slope|
+-----+-----+-----+-----+
|Brazil |2020-01-09 00:00:00|null  |[]
|null |
|Brazil |2020-01-23 00:00:00|0      |[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
|null |
|Brazil |2020-02-06 00:00:00|0      |[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
|null |
|Brazil |2020-02-20 00:00:00|4      |[0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 2]
|null |
|Brazil |2020-03-05 00:00:00|368    |[0, 9, 0, 7, 5, 6, 7, 14, 99, 0, 11, 38, 121, 5
1]
|81.2 |
|Brazil |2020-03-19 00:00:00|6464   |[249, 172, 228, 525, 378, 323, 307, 431, 432, 4
87, 352, 323, 1138, 1119]
|88.8 |
|Brazil |2020-04-02 00:00:00|21484  |[1208, 1012, 1304, 770, 1031, 1873, 2136, 1922,
1546, 1089, 1465, 1238, 1832, 3058]
|89.3 |
|Brazil |2020-04-16 00:00:00|51365  |[2105, 3257, 2976, 1996, 2089, 2336, 2678, 4279
, 4007, 5281, 3776, 4346, 5789, 6450]
|89.8 |
|Brazil |2020-04-30 00:00:00|110452|[7502, 5015, 4898, 4726, 6794, 6835, 11156, 916
2, 11121, 9167, 6638, 6895, 8620, 11923]
|89.8 |
|Brazil |2020-05-14 00:00:00|81748 |[13028, 17126, 13220, 7569, 14288, 16517]
|89.4 |
+-----+-----+-----+-----+
```

```
# Adding the slop sign ( increase/decrease in binned newConfirmed growth)
```

```
# groupin by cantry and sign then counting
```

```
dfc8 = dfc7\
    .withColumn( "slopeSign", fun.when( dfc7.slope<0., "-" ).otherwise("+") )\
    .groupBy( "country", "slopeSign" )\
    .agg( fun.count(dfc7.slope).alias("count") )\
    .orderBy("country")
# .groupBy("country")\
# .agg( fun.collect_list("count").alias("count-+") )\
```

```
dfc8.show(500)
dfc8.filter( "country = 'Norway'" ).show()
```

country	slopeSign	count
Afghanistan	+	6
Albania	+	1
Albania	-	5
Algeria	-	2
Algeria	+	4
Andorra	-	2
Andorra	+	2
Angola	-	1
Angola	+	1
Antigua and Barbuda	+	0
Argentina	+	6
Armenia	+	5
Armenia	-	1
Australia	-	4
Australia	+	2
Austria	-	4
Austria	+	3
Azerbaijan	+	3

```
# considering the increase numbers
```

```
dfc8.filter("slopeSign='+'").orderBy("count", "country", ascending=False).show(200)
```

country	slopeSign	count
Singapore	+	6
Saudi Arabia	+	6
Pakistan	+	6
Mexico	+	6
Kuwait	+	6
Iraq	+	6
India	+	6
Egypt	+	6
Colombia	+	6
Chile	+	6
Brazil	+	6
Argentina	+	6
Afghanistan	+	6
United Arab Emirates	+	5
South Africa	+	5

	Russia	+	5
	Qatar	+	5

this is compatible for Brazil plot (binned newConfimd vs. date). Explaine why for the last bin we observe a decreas?

